

ass1

April 23, 2024

```
[71]: #shree swami samarth
```

```
[1]: import pandas as pd
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[2]: from keras.models import Sequential
from keras.layers import Dense, Activation
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import metrics
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
```

```
[3]: import warnings
warnings.filterwarnings("ignore")
```

```
[4]: df = pd.read_csv("boston.csv")
```

```
[5]: df
```

```
[5]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	ptratio	b	lstat	MEDV
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2
..
501	21.0	391.99	9.67	22.4
502	21.0	396.90	9.08	20.6
503	21.0	396.90	5.64	23.9
504	21.0	393.45	6.48	22.0
505	21.0	396.90	7.88	11.9

[506 rows x 14 columns]

```
[6]: df.head()
```

```
[6]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
[7]: df.describe()
```

```
[7]:
```

	crim	zn	indus	chas	nox	rm	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	

	age	dis	rad	tax	ptratio	b	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	

min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	lstat	MEDV
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
[8]: df.tail()
```

```
[8]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	

	b	lstat	MEDV
501	391.99	9.67	22.4
502	396.90	9.08	20.6
503	396.90	5.64	23.9
504	393.45	6.48	22.0
505	396.90	7.88	11.9

```
[9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   crim        506 non-null    float64
1   zn          506 non-null    float64
2   indus       506 non-null    float64
3   chas        506 non-null    int64
4   nox         506 non-null    float64
5   rm          506 non-null    float64
6   age         506 non-null    float64
7   dis         506 non-null    float64
```

```

8   rad      506 non-null    int64
9   tax      506 non-null    int64
10  ptratio  506 non-null    float64
11  b        506 non-null    float64
12  lstat    506 non-null    float64
13  MEDV     506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB

```

```

[10]: # Separate the features (X) and target variable (y)
#X = data.drop('target_column', axis=1)
#y = data['target_column']

X = df.drop('MEDV', axis=1)

Y = df['MEDV']

```

```

[11]: X

```

```

[11]:      crim    zn  indus  chas   nox    rm   age   dis  rad  tax  \
0    0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900   1  296
1    0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671   2  242
2    0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671   2  242
3    0.03237   0.0   2.18    0  0.458  6.998  45.8  6.0622   3  222
4    0.06905   0.0   2.18    0  0.458  7.147  54.2  6.0622   3  222
..      ...   ...   ...   ...   ...   ...   ...   ...   ...
501  0.06263   0.0  11.93    0  0.573  6.593  69.1  2.4786   1  273
502  0.04527   0.0  11.93    0  0.573  6.120  76.7  2.2875   1  273
503  0.06076   0.0  11.93    0  0.573  6.976  91.0  2.1675   1  273
504  0.10959   0.0  11.93    0  0.573  6.794  89.3  2.3889   1  273
505  0.04741   0.0  11.93    0  0.573  6.030  80.8  2.5050   1  273

      ptratio    b  lstat
0         15.3  396.90   4.98
1         17.8  396.90   9.14
2         17.8  392.83   4.03
3         18.7  394.63   2.94
4         18.7  396.90   5.33
..      ...   ...   ...
501        21.0  391.99   9.67
502        21.0  396.90   9.08
503        21.0  396.90   5.64
504        21.0  393.45   6.48
505        21.0  396.90   7.88

[506 rows x 13 columns]

```

```
[12]: Y
```

```
[12]: 0      24.0
      1      21.6
      2      34.7
      3      33.4
      4      36.2
      ...
      501    22.4
      502    20.6
      503    23.9
      504    22.0
      505    11.9
      Name: MEDV, Length: 506, dtype: float64
```

```
[13]: #This part scales the input features and target variables using the
      ↪StandardScaler from scikit-learn. The scaler is fit on the training data,
      ↪and then both training and test data are transformed using the same scaler.
      scaler = StandardScaler().fit(X)
      X = scaler.transform(X)
```

```
[14]: # Split the data into train and test sets
      # testsize 0.2 Percentage of data to use for testing
      x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2,
      ↪random_state=42)
```

```
[15]: x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
[15]: ((404, 13), (404,), (102, 13), (102,))
```

```
[16]: from keras.layers import Input

      model = Sequential()
      #model.add for 7 time
      model.add(Input(shape = (x_train.shape[1],))) #input layer with input shape
      model.add(Dense(12, activation = 'relu'))
      model.add(Dense(13, activation = 'relu'))
      model.add(Dense(19, activation = 'relu'))
      model.add(Dense(19, activation = 'relu'))
      model.add(Dense(19, activation = 'relu'))
      model.add(Dense(1, activation = 'linear'))
```

```
[17]: model.summary()
```

```
Model: "sequential"
```

Layer (type) ↳Param #	Output Shape	
dense (Dense) ↳168	(None, 12)	↳
dense_1 (Dense) ↳169	(None, 13)	↳
dense_2 (Dense) ↳266	(None, 19)	↳
dense_3 (Dense) ↳380	(None, 19)	↳
dense_4 (Dense) ↳380	(None, 19)	↳
dense_5 (Dense) ↳ 20	(None, 1)	↳

Total params: 1,383 (5.40 KB)

Trainable params: 1,383 (5.40 KB)

Non-trainable params: 0 (0.00 B)

```
[18]: #The model is compiled with the Adam optimizer, mean squared error (MSE) as the
      ↳loss function, and mean absolute error (MAE) as an additional evaluation
      ↳metric
```

```
model.compile(loss='mean_squared_error', optimizer='adam' )
```

```
[19]: #The model is trained on the scaled training data (x_train and y_train) for 20
      ↳epochs with a batch size of 32. The scaled test data (x_test and y_test) is
      ↳used for validation during training. The training history, including loss
      ↳and metrics values for each epoch, is stored in the history variable.
      # Fitting the data to the model
```

```
history = model.fit(x_train, y_train, epochs=100, batch_size=16, verbose = 0)
```

```
[20]: mse = model.evaluate(x_test, y_test, verbose = 0)
      mse
```

[20]: 12.522753715515137

```
[21]: #After training, make prediction on test data
y_pred = model.predict(x_test)
```

4/4 0s 16ms/step

```
[22]: #calculate RMSE
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error(RMSE):" , rmse)
```

Root Mean Squared Error(RMSE): 3.538750386302523

```
[23]: y_pred = model.predict(x_test)
y_pred
```

4/4 0s 5ms/step

```
[23]: array([[27.454987],
 [32.89278 ],
 [17.452944],
 [25.303612],
 [15.917062],
 [19.693556],
 [17.11813 ],
 [14.602602],
 [25.2784  ],
 [19.036325],
 [21.321346],
 [17.513882],
 [ 6.464231],
 [17.906958],
 [16.52573 ],
 [22.423206],
 [20.137053],
 [10.732348],
 [47.894962],
 [14.764127],
 [24.974274],
 [27.381502],
 [14.572911],
 [23.990072],
 [18.044992],
 [19.46825 ],
 [20.81469 ],
 [13.034436],
 [19.293957],
 [17.511202],
```

[24.493526],
[23.551088],
[19.056452],
[24.560196],
[15.956036],
[15.600607],
[33.454098],
[20.494028],
[18.553719],
[25.743597],
[14.882683],
[31.436174],
[51.556454],
[18.139557],
[25.907238],
[18.20188],
[15.068213],
[26.377502],
[19.662882],
[25.35477],
[18.736301],
[35.250175],
[16.665594],
[24.05209],
[40.416695],
[22.207521],
[17.158577],
[33.774757],
[23.876213],
[17.610085],
[24.261],
[33.345276],
[31.48246],
[19.758612],
[23.175898],
[17.36896],
[16.825018],
[24.317253],
[29.672707],
[12.728228],
[21.778105],
[28.38147],
[8.041182],
[21.330404],
[20.856318],
[6.80382],
[19.844055],


```
[48.756413],
[11.667135],
[12.927367],
[18.647955],
[13.615104],
[19.72286 ],
[12.643882],
[19.615639],
[26.34903 ],
[15.533831],
[24.353762],
[25.85339 ],
[18.436605],
[24.297964],
[ 6.671335],
[19.005749],
[19.093147],
[34.084385],
[20.42649 ],
[28.166824],
[ 6.129143],
[12.479241],
[13.914244],
[22.049635],
[21.373142]], dtype=float32)
```

```
[24]: y_test_np = np.array(y_test)
y_pred_np = np.array(y_pred)

# Flatten or ravel the arrays
y_test_flat = y_test_np.flatten()
y_pred_flat = y_pred_np.flatten()

# Create the DataFrame
results_df = pd.DataFrame({'Actual Values': y_test_flat, 'Predicted Values':
    ↪ y_pred_flat})

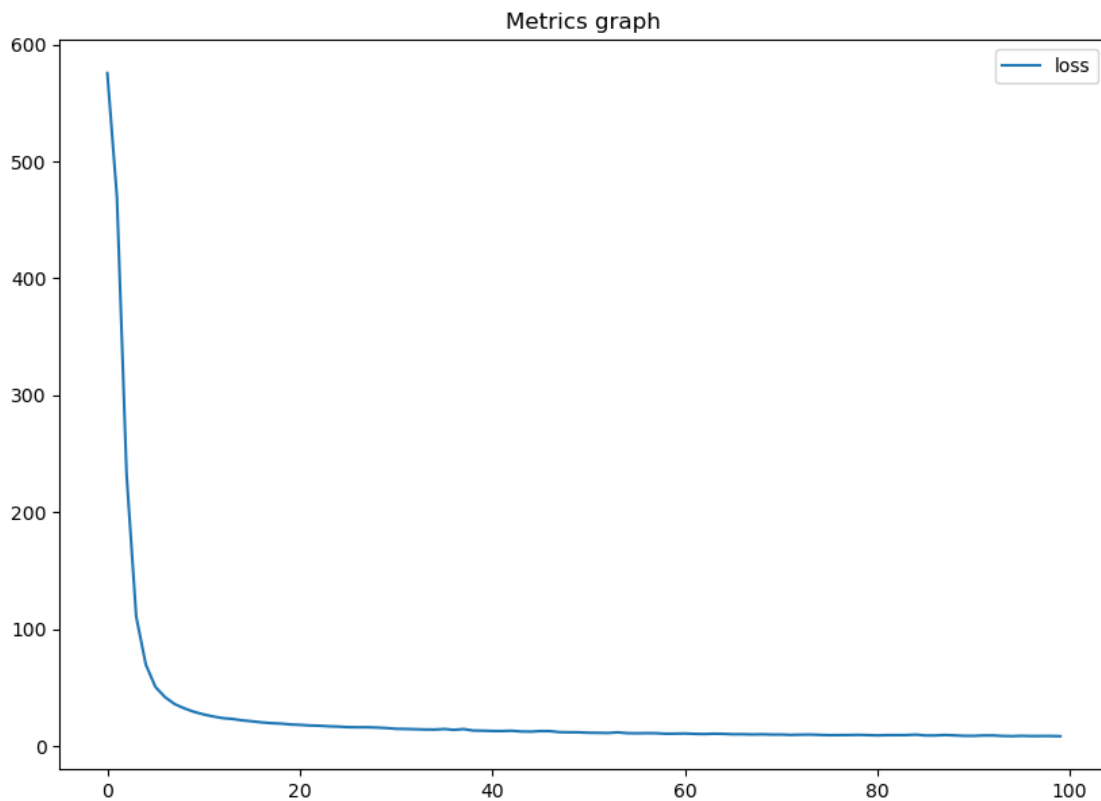
# Print the DataFrame
print(results_df)
```

	Actual Values	Predicted Values
0	23.6	27.454987
1	32.4	32.892780
2	13.6	17.452944
3	22.8	25.303612
4	16.1	15.917062
..


97	17.9	6.129143
98	9.6	12.479241
99	17.2	13.914244
100	22.5	22.049635
101	21.4	21.373142

[102 rows x 2 columns]

```
[25]: # This part creates a line plot using the training history stored in history.
      ↪ history. The plot displays the training and validation loss and MAE values
      ↪ over the epochs.
      pd.DataFrame(history.history).plot(figsize=(10,7))
      plt.title("Metrics graph")
      plt.show()
```



```
[26]: # plotting the true test values (y_test) against the predicted values (y_pred).
      ↪ A regression line is also plotted to visualize the relationship between the
      ↪ true and predicted values.
      sns.regplot(x=y_test, y=y_pred)
      plt.title("Regression Line for Predicted values")
      plt.show()
```



[27]: *# Evaluating <keras.src.callbacks.History at 0x7c5098f9db40> the model*

```
def regression_metrics_display(y_test, y_pred):
    print(f"MAE is {metrics.mean_absolute_error(y_test, y_pred)}")
    print(f"MSE is {metrics.mean_squared_error(y_test, y_pred)}")
    print(f"R2 score is {metrics.r2_score(y_test, y_pred)}")
```

[28]: *#fun call*

```
regression_metrics_display(y_test, y_pred)
```

MAE is 2.2527441146327

MSE is 12.522754296556258

R2 score is 0.8292362161491259

[29]: *#basic*

```
df.shape
```

[29]: (506, 14)

[30]: df.dtypes

```
[30]: crim      float64
      zn        float64
      indus     float64
      chas      int64
      nox       float64
      rm        float64
      age       float64
      dis       float64
      rad       int64
      tax       int64
      ptratio   float64
      b         float64
      lstat     float64
      MEDV      float64
      dtype: object
```

```
[31]: df.isna().sum()
```

```
[31]: crim      0
      zn        0
      indus     0
      chas      0
      nox       0
      rm        0
      age       0
      dis       0
      rad       0
      tax       0
      ptratio   0
      b         0
      lstat     0
      MEDV      0
      dtype: int64
```

```
[32]: df.describe()
```

```
[32]:
```

	crim	zn	indus	chas	nox	rm \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

	age	dis	rad	tax	ptratio	b \
--	-----	-----	-----	-----	---------	-----


count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	lstat	MEDV
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000

```
[33]: #data visualization
import seaborn as sns

sns.displot(df.MEDV)
```

```
[33]: <seaborn.axisgrid.FacetGrid at 0x29cda3005d0>
```



```
[34]: import seaborn as sns
import matplotlib.pyplot as plt
# Assuming 'df' is your DataFrame
correlation = df.corr()

# Creating a heatmap with correlation values annotated
fig, axes = plt.subplots(figsize=(15, 12))
sns.heatmap(correlation, square=True, annot=True)

# Adding title
plt.title('Correlation Matrix Heatmap')

# Showing the plot
plt.show()
```



[]:

[]:

[]: